

Thus the LeNet architecture has been implemented successfully

## **AD18511 – DEEP LEARNING LABORATORY**

**DATE:**

**EX.NO: 9**

### **IMPLEMENTATION OF ALEXNET ARCHITECTURE.**

#### **AIM:**

To implement alexnet architecture for image recognition and classification tasks.

#### **DESCRIPTION:**

- The AlexNet CNN architecture won the 2012 ImageNet ILSVRC challenges of deep learning algorithm by a large variance by achieving 17% with top-5 error rate as the second best achieved 26%!
- It was introduced by Alex Krizhevsky (name of founder), The Ilya Sutskever and Geoffrey Hinton are quite similar to LeNet-5, only much bigger and deeper and it was introduced first to stack convolutional layers directly on top of each other models, instead of stacking a pooling layer top of each on CN network convolutional layer.
- AlexNNet has 60 million parameters as AlexNet has total 8 layers, 5 convolutional and 3 fully connected layers.
- AlexNNet is first to execute (ReLU) Rectified Linear Units as activation functions.
- it was the first CNN architecture that uses GPU to improve the performance.

#### **PROGRAM:**

```
# Importing Libraries
import tensorflow as tf
from tensorflow import keras
from keras.layers import Conv2D, MaxPooling2D, Dense, Flatten, Dropout, BatchNormalization
from tensorflow.keras.optimizers import Adam,SGD
from tensorflow.keras.callbacks import TensorBoard

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split

import os
import time

train_df = pd.read_csv('/fashion-mnist_train.csv',sep=',')
test_df = pd.read_csv('/fashion-mnist_test.csv', sep = ',')

from google.colab import drive
drive.mount('/content/drive')

train.head(10)
train_data= np.array(train_df, dtype= 'float32')
test_data= np.array(test_df, dtype= 'float32')

x_train = train_data[:,1:]/255
y_train = train_data[:,0]
x_test= test_data[:,1:]/255
y_test=test_data[:,0]
```

```
# Example of training label content
print(y_train[0], y_train[43], y_train[1923])
print("Minimum value of training labels", y_train.min())
print("Maximum value of training labels", y_train.max())
```

### OUTPUT:

```
0.0 6.0 2.0
Minimum value of training labels 0.0
Maximum value of training labels 9.0
```

```
x_train,x_validate,y_train,y_validate = train_test_split(x_train,y_train,test_size = 0.2,random_state = 12345)
x_train.shape, x_validate.shape,y_train.shape,y_validate.shape
```

### OUTPUT:

```
((8000, 784), (2000, 784), (8000,), (2000,))
```

```
class_names= ['Tshirt', 'Trouser', 'Pullover', 'Dress', 'Coat',
              'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle boot']
```

```
plt.figure(figsize=(10,10))
for i in range(36):
    plt.subplot(6,6, i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(x_train[i].reshape((28,28)))
    label_index= int(y_train[i])
    plt.title(class_names[label_index])
plt.show()
```

### OUTPUT:



```

image_rows = 28
image_cols = 28
batch_size = 4096
image_shape = (image_rows,image_cols,1)

x_train= x_train.reshape(x_train.shape[0], *image_shape)
x_test = x_test.reshape(x_test.shape[0],*image_shape)
x_validate = x_validate.reshape(x_validate.shape[0],*image_shape)
x_train.shape, x_test.shape, x_validate.shape

# Alexnet
model= tf.keras.Sequential([Conv2D(filters=96, kernel_size=(11,11), strides=(4,4), activation= 'relu', input_shape=
image_shape),
                           BatchNormalization(),
                           MaxPooling2D(pool_size=(3,3), strides=(2,2), padding='same'),
                           Conv2D(filters=256, kernel_size=(5,5), strides=(1,1), activation='relu', padding='same'),
                           BatchNormalization(),
                           MaxPooling2D(pool_size=(3,3), strides=(2,2), padding='same'),
                           Conv2D(filters=384, kernel_size=(3,3), strides=(1,1), activation='relu', padding='same'),
                           BatchNormalization(),
                           Conv2D(filters=384, kernel_size=(3,3), strides=(1,1), activation='relu', padding='same'),
                           BatchNormalization(),
                           Conv2D(filters=256, kernel_size=(3,3), strides=(1,1), activation='relu', padding='same'),
                           BatchNormalization(),
                           MaxPooling2D(pool_size=(3,3), strides=(2,2), padding='same'),
                           Flatten(),
                           Dense(4096, activation='relu'),
                           Dropout(0.5),
                           Dense(4096, activation='relu'),
                           Dropout(0.5),
                           Dense(10, activation='softmax')
                           ])

model.summary()

```

## OUTPUT:

Model: "sequential"

| Layer (type)                                | Output Shape      | Param # |
|---|-------------------|---------|
| conv2d (Conv2D)                             | (None, 5, 5, 96)  | 11712   |
| batch_normalization (Batch Normalization)   | (None, 5, 5, 96)  | 384     |
| max_pooling2d (MaxPooling2D)                | (None, 3, 3, 96)  | 0       |
| conv2d_1 (Conv2D)                           | (None, 3, 3, 256) | 614656  |
| batch_normalization_1 (Batch Normalization) | (None, 3, 3, 256) | 1024    |

max\_pooling2d\_1 (MaxPoolin (None, 2, 2, 256) 0

g2D)

|   |                   |         |
|---|-------------------|---------|
| conv2d_2 (Conv2D)                           | (None, 2, 2, 384) | 885120  |
| batch_normalization_2 (Batch Normalization) | (None, 2, 2, 384) | 1536    |
| conv2d_3 (Conv2D)                           | (None, 2, 2, 384) | 1327488 |
| batch_normalization_3 (Batch Normalization) | (None, 2, 2, 384) | 1536    |
| conv2d_4 (Conv2D)                           | (None, 2, 2, 256) | 884992  |
| batch_normalization_4 (Batch Normalization) | (None, 2, 2, 256) | 1024    |
| max_pooling2d_2 (MaxPooling2D)              | (None, 1, 1, 256) | 0       |

|                     |              |          |
|---------------------|--------------|----------|
| flatten (Flatten)   | (None, 256)  | 0        |
| dense (Dense)       | (None, 4096) | 1052672  |
| dropout (Dropout)   | (None, 4096) | 0        |
| dense_1 (Dense)     | (None, 4096) | 16781312 |
| dropout_1 (Dropout) | (None, 4096) | 0        |
| dense_2 (Dense)     | (None, 10)   | 40970    |

---

Total params: 21604426 (82.41 MB)  
Trainable params: 21601674 (82.40 MB)  
Non-trainable params: 2752 (10.75 KB)

---

```
from tensorflow.keras.optimizers import SGD
```

```
def lr_schedule(epoch):  
    lr = 0.01  
    if epoch > 50:  
        lr *= 0.1  
    elif epoch > 75:  
        lr *= 0.01  
    return lr
```

```
sgd = SGD(learning_rate=0.01, momentum=0.9, nesterov=True)  
model.compile(loss="sparse_categorical_crossentropy", optimizer=sgd, metrics=["accuracy"])  
from tensorflow.keras.callbacks import LearningRateScheduler  
lr_scheduler = LearningRateScheduler(lr_schedule)
```

# train!

```
early_stopping_cb = keras.callbacks.EarlyStopping(monitor='val_loss', patience=30, verbose=1, mode='min')
history= model.fit(x_train, y_train, epochs=5, verbose=1, callbacks=[early_stopping_cb],
                    validation_data=(x_validate, y_validate))
```

## OUTPUT:

Epoch 1/5

250/250 [=====] - 220s 881ms/step - loss: 0.4928 - accuracy: 0.8270 - val\_loss: 0.4818 - val\_accuracy: 0.8220

Epoch 2/5

250/250 [=====] - 218s 872ms/step - loss: 0.4541 - accuracy: 0.8384 - val\_loss: 0.4485 - val\_accuracy: 0.8420

Epoch 3/5

250/250 [=====] - 213s 852ms/step - loss: 0.4073 - accuracy: 0.8497 - val\_loss: 0.4617 - val\_accuracy: 0.8310

Epoch 4/5

250/250 [=====] - 216s 863ms/step - loss: 0.4110 - accuracy: 0.8534 - val\_loss: 0.4099 - val\_accuracy: 0.8565


Epoch 5/5

250/250 [=====] - 222s 890ms/step - loss: 0.3683 - accuracy: 0.8661 - val\_loss: 0.4187 - val\_accuracy: 0.8505

```
plt.figure(figsize=(10,10))
plt.subplot(2,2,1)
plt.plot(history.history['loss'], label='loss',color='r')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.legend()
plt.title('Training - Loss Function')
```

```
plt.subplot(2, 2, 2)
plt.plot(history.history['accuracy'], label='Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.legend()
plt.title('Train - Accuracy')
```

## OUTPUT:



```
model_evaluation_results = model.evaluate(x_test, y_test, batch_size=32, verbose=2)
print("The test loss is", model_evaluation_results[0])
print("The test accuracy is", model_evaluation_results[1])
```

### OUTPUT:

```
533/533 - 42s - loss: nan - accuracy: 0.8438 - 42s/epoch - 79ms/step
The test loss is nan
The test accuracy is 0.8438380360603333
```

```
# Prediction on test images using model.predict() method
practical_test_images = x_test[:10]
prediction_probabilites = model.predict(practical_test_images)
prediction_probabilites[:3]
```

### OUTPUT:

```
1/1 [=====
=====
=====
1.62392721e-01, 6.00572117e-03, 3.07703376e-01, 3.86583386e-03,
8.07231665e-02, 4.55609756e-03],
[2.13454390e-04, 1.31126042e-04, 2.42721246e-04, 1.57689094e-04,
2.41081478e-04, 5.40605932e-03, 3.75541451e-04, 3.94529492e-01,
3.35998135e-04, 5.983666857e-01],
[1.77102792e-03, 9.13637981e-04, 1.02652036e-01, 3.12519167e-03,
5.54388940e-01, 6.66437962e-04, 3.32357943e-01, 6.95253024e-04,
2.90616718e-03, 5.23336872e-04], dtype=float32)
```

```
# Clean up model prediction using argmax to find the index of the largest probability
def derive_predicted_classes(prediction_probabilites):
    batch_prediction = []
    for vector in prediction_probabilites:
        batch_prediction.append(np.argmax(vector))
    return batch_prediction

model_prediction = derive_predicted_classes(prediction_probabilites)
model_prediction
```


### OUTPUT:

```
[6, 9, 4, 0, 3, 6, 4, 5, 4, 8]
```

```
# Visualise the prediction result
plt.figure(figsize=(10,10))
for i in range(len(practical_test_images)):
    plt.subplot(5,5, i+1)
    plt.axis("off")
    plt.grid(False)
    plt.imshow(practical_test_images[i])
    plt.title(class_names[model_prediction[i]])
plt.show()
```



## OUTPUT:



```
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
# Reshape x_test to have the correct shape
```

```
x_test = x_test.reshape(x_test.shape[0], image_rows, image_cols, 1)
```

```
# Calculate accuracy score
```

```
y_pred = model.predict(x_test)
```

```
y_pred_classes = np.argmax(y_pred, axis=1)
```

```
accuracy = accuracy_score(y_test, y_pred_classes)
```

```
print("Accuracy:", accuracy)
```

## OUTPUT:

533/533 [=====] - 41s 78ms/step Accuracy: 0.8438380281690141

```
# Compute confusion matrix
```

```
confusion_mtx = confusion_matrix(y_test, y_pred_classes)
```

```
# Display confusion matrix as a heatmap
```

```
import seaborn as sns
```

```
plt.figure(figsize=(10,8))
```

```
sns.heatmap(confusion_mtx, annot=True, fmt='d', cmap='Blues', xticklabels=class_names, yticklabels=class_names)
```


```
plt.xlabel('Predicted Labels')
```

```
plt.ylabel('True Labels')
```

```
plt.title('Confusion Matrix')
```

```
plt.show()
```

## OUTPUT:



```
# Generate and display classification report
classification_rep = classification_report(y_test, y_pred_classes,
target_names=class_names) print("Classification Report:\n", classification_rep)
```

## OUTPUT:

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| Tshirt       | 0.81      | 0.81   | 0.81     | 1770    |
| Trouser      | 0.99      | 0.94   | 0.97     | 1700    |
| Pullover     | 0.71      | 0.79   | 0.75     | 1677    |
| Dress        | 0.86      | 0.89   | 0.87     | 1725    |
| Coat         | 0.72      | 0.81   | 0.76     | 1639    |
| Sandal       | 0.97      | 0.86   | 0.91     | 1695    |
| Shirt        | 0.63      | 0.51   | 0.56     | 1704    |
| Sneaker      | 0.83      | 0.98   | 0.90     | 1761    |
| Bag          | 0.97      | 0.95   | 0.96     | 1675    |
| Ankle boot   | 0.97      | 0.91   | 0.94     | 1694    |
| accuracy     |           |        | 0.84     | 17040   |
| macro avg    | 0.85      | 0.84   | 0.84     | 17040   |
| weighted avg | 0.85      | 0.84   | 0.84     | 17040   |

## RESULT:

Thus the ALEXNet architecture has been implemented successfully.