

## **AD18511 – DEEP LEARNING LABORATORY**

**DATE:**

**EX.NO: 10**

### **IMPLEMENTATION OF VGG ARCHITECTURE.**

**AIM:**

To use the VGG (Visual Geometric group) architecture to CIFAR-10 datasets using Tensorflow.

**DESCRIPTION:**

- VGG stands for Visual Geometry Group; it is a standard deep Convolutional Neural Network (CNN) architecture with multiple layers.
- The —deep refers to the number of layers with VGG-16 or VGG-19 consisting of 16 and 19 convolutional layers.
- The VGG architecture is the basis of ground-breaking object recognition models. Developed as a deep neural network, the VGGNet also surpasses baselines on many tasks and datasets beyond ImageNet. Moreover, it is now still one of the most popular image recognition architectures.
- The VGG network is constructed with very small convolutional filters. The VGG-16 consists of 13 convolutional layers and three fully connected layers.
- Input: The VGGNet takes in an image input size of  $224 \times 224$
- Convolutional Layers: VGG's convolutional layers leverage a minimal receptive field, i.e.,  $3 \times 3$ , the smallest possible size that still captures up/down and left/right.
- Hidden Layers: All the hidden layers in the VGG network use ReLU. VGG does not usually leverage Local Response Normalization (LRN) as it increases memory consumption and training time.
- Fully-Connected Layers: The VGGNet has three fully connected layers. Out of the three layers, the first two have 4096 channels each, and the third has 1000 channels, 1 for each class.

**PROGRAM:**

```
import tensorflow as tf
from tensorflow.keras import datasets, layers, models
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, classification_report
import numpy as np

# Load and preprocess CIFAR-10 dataset
(train_images, train_labels), (test_images, test_labels) = datasets.cifar10.load_data()
train_images, test_images = train_images / 255.0, test_images / 255.0
```

**OUTPUT:**

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>  
170498071/170498071 [=====] - 14s 0us/step

```
model = models.Sequential()
model.add(layers.Conv2D(64, (3, 3), activation='relu', input_shape=(32, 32, 3), padding='same'))
model.add(layers.BatchNormalization(scale=False)) # Add BatchNormalization layer
model.add(layers.Conv2D(64, (3, 3), activation='relu', padding='same'))
model.add(layers.MaxPooling2D((2, 2)))
```

```

model.add(layers.Conv2D(128, (3, 3), activation='relu', input_shape=(32, 32, 3), padding='same'))
model.add(layers.BatchNormalization(scale=False)) # Add BatchNormalization layer
model.add(layers.Conv2D(128, (3, 3), activation='relu', padding='same'))
model.add(layers.MaxPooling2D((2, 2)))

model.add(layers.Conv2D(256, (3, 3), activation='relu', input_shape=(32, 32, 3), padding='same'))
model.add(layers.BatchNormalization(scale=False)) # Add BatchNormalization layer
model.add(layers.Conv2D(256, (3, 3), activation='relu', padding='same'))
model.add(layers.MaxPooling2D((2, 2)))

model.add(layers.Conv2D(512, (3, 3), activation='relu', input_shape=(32, 32, 3), padding='same'))
model.add(layers.BatchNormalization(scale=False)) # Add BatchNormalization layer
model.add(layers.Conv2D(512, (3, 3), activation='relu', padding='same'))
model.add(layers.MaxPooling2D((2, 2)))

model.add(layers.Conv2D(512, (3, 3), activation='relu', input_shape=(32, 32, 3), padding='same'))
model.add(layers.BatchNormalization(scale=False)) # Add BatchNormalization layer
model.add(layers.Conv2D(512, (3, 3), activation='relu', padding='same'))
model.add(layers.MaxPooling2D((2, 2)))

model.add(layers.Flatten())
model.add(layers.Dense(4096, activation='relu'))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(4096, activation='relu'))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(10, activation='softmax'))

# creating the model
model.summary()

```

## OUTPUT:

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 64)	1792
batch_normalization (Batch Normalization)	(None, 32, 32, 64)	192
conv2d_1 (Conv2D)	(None, 32, 32, 64)	36928
max_pooling2d (MaxPooling2D)	(None, 16, 16, 64)	0
conv2d_2 (Conv2D)	(None, 16, 16, 128)	73856
batch_normalization_1 (Batch Normalization)	(None, 16, 16, 128)	384
conv2d_3 (Conv2D)	(None, 16, 16, 128)	147584

max_pooling2d_1 (MaxPoolin g2D)	(None, 8, 8, 128)	0
conv2d_4 (Conv2D)	(None, 8, 8, 256)	295168
batch_normalization_2 (Bat chNormalization)	(None, 8, 8, 256)	768
conv2d_5 (Conv2D)	(None, 8, 8, 256)	590080
max_pooling2d_2 (MaxPoolin g2D)	(None, 4, 4, 256)	0
conv2d_6 (Conv2D)	(None, 4, 4, 512)	1180160
batch_normalization_3 (Bat chNormalization)	(None, 4, 4, 512)	1536
conv2d_7 (Conv2D)	(None, 4, 4, 512)	2359808
max_pooling2d_3 (MaxPoolin g2D)	(None, 2, 2, 512)	0
conv2d_8 (Conv2D)	(None, 2, 2, 512)	2359808
batch_normalization_4 (Bat chNormalization)	(None, 2, 2, 512)	1536
conv2d_9 (Conv2D)	(None, 2, 2, 512)	2359808
max_pooling2d_4 (MaxPoolin g2D)	(None, 1, 1, 512)	0
flatten (Flatten)	(None, 512)	0
dense (Dense)	(None, 4096)	2101248
dropout (Dropout)	(None, 4096)	0
dense_1 (Dense)	(None, 4096)	16781312
dropout_1 (Dropout)	(None, 4096)	0
dense_2 (Dense)	(None, 10)	40970

---

Total params: 28332938 (108.08 MB)  
 Trainable params: 28329994 (108.07 MB)  
 Non-trainable params: 2944 (11.50 KB)

---

```
model.compile(optimizer='adam',  
              loss='sparse_categorical_crossentropy',  
              metrics=['accuracy'])
```

```
# Train the model
```

```
history = model.fit(train_images, train_labels, epochs=10,  
                    validation_data=(test_images, test_labels))
```

## OUTPUT:

Epoch 1/10

1563/1563 [=====] - 57s 27ms/step - loss: 1.6695 - accuracy: 0.3731 - val\_loss: 1.6695 - val\_accuracy: 0.4469

Epoch 2/10

1563/1563 [=====] - 39s 25ms/step - loss: 1.1939 - accuracy: 0.5878 - val\_loss: 1.3521 - val\_accuracy: 0.5481

Epoch 3/10

1563/1563 [=====] - 40s 25ms/step - loss: 0.9247 - accuracy: 0.6933 - val\_loss: 0.8926 - val\_accuracy: 0.6929

Epoch 4/10

1563/1563 [=====] - 40s 25ms/step - loss: 0.7902 - accuracy: 0.7438 - val\_loss: 0.8881 - val\_accuracy: 0.7111

Epoch 5/10

1563/1563 [=====] - 40s 25ms/step - loss: 0.6856 - accuracy: 0.7808 - val\_loss: 0.8698 - val\_accuracy: 0.7427

Epoch 6/10

1563/1563 [=====] - 40s 25ms/step - loss: 0.6300 - accuracy: 0.8035 - val\_loss: 0.8138 - val\_accuracy: 0.7857

Epoch 7/10

1563/1563 [=====] - 41s 26ms/step - loss: 0.5603 - accuracy: 0.8251 - val\_loss: 1.0362 - val\_accuracy: 0.7885

Epoch 8/10

1563/1563 [=====] - 45s 29ms/step - loss: 0.4660 - accuracy: 0.8551 - val\_loss: 0.7054 - val\_accuracy: 0.7873

Epoch 9/10


1563/1563 [=====] - 39s 25ms/step - loss: 0.4313 - accuracy: 0.8671 - val\_loss: 2.1532 - val\_accuracy: 0.7867

Epoch 10/10

1563/1563 [=====] - 39s 25ms/step - loss: 0.3916 - accuracy: 0.8817 - val\_loss: 4.2583 - val\_accuracy: 0.8076

```
plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label='val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.ylim([0, 1])
plt.legend(loc='lower right')
plt.show()
```

## OUTPUT:



```
test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)
print(f"Test accuracy: {test_acc}")
```

## OUTPUT:

313/313 - 2s - loss: 4.2583 - accuracy: 0.8076 - 2s/epoch - 6ms/step Test accuracy: 0.8076000213623047

```
# Make predictions
predictions = np.argmax(model.predict(test_images), axis=-1)
```

## OUTPUT:

313/313 [=====] - 2s 6ms/step

```
# Make predictions
predictions = np.argmax(model.predict(test_images), axis=-1)
# Generate classification report
report = classification_report(test_labels, predictions)
```

```
# Print confusion matrix and classification report
print("Confusion Matrix:\n", cm)
print("Classification Report:\n", report)
```

## OUTPUT:

Confusion Matrix:

```
[[825 8 43 37 10 3 7 26 6 35]
 [ 7 923 8 3 2 0 5 3 1 48]
 [ 46  0 655 31 76 81 56 43 0 12]
 [  4  3  29 549 35 258 62 44 2 14]
 [  4  1  16 20 795 39 25 97 1 2]
 [  1  1  9 73 19 838 7 50 0 2]
 [  2  1 28 33 19 16 886 10 2 3]
 [  4  1  4 12 14 26 1 931 0 7]
 [ 79 52 10 11 4 2 17 10 748 67]
 [  8 26 2 5 1 3 4 14 11 926]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.84	0.82	0.83	1000
1	0.91	0.92	0.92	1000
2	0.81	0.66	0.73	1000
3	0.71	0.55	0.62	1000
4	0.82	0.80	0.81	1000
5	0.66	0.84	0.74	1000
6	0.83	0.89	0.86	1000
7	0.76	0.93	0.84	1000
8	0.97	0.75	0.84	1000
9	0.83	0.93	0.88	1000
accuracy		0.81		10000
macro avg	0.81	0.81	0.81	10000
weighted avg	0.81	0.81	0.81	10000

## RESULT:

Thus the VGG architecture has been implemented successfully.